

Using Python to Speed-up Applications with GPUs

Continuum Analytics

Travis Oliphant & Siu Kwan Lam

Agenda

- Introduction
 - Python
 - Numba
 - NumbaPro
- Lab 1
- Lab 2

Why Python?

Pros

- A dynamic scripting language
- Rapid development
- Rich libraries — NumPy, SciPy
- Great glue language to connect native libraries

Cons

- Hard to parallelize because of the GIL
- Slow execution speed

Our Solution: Speedup Python with Numba

Numba

- An opensource **JIT** compiler for **array-oriented programming** in CPython
- Turn numerical loops into fast native code
- Maximize hardware utilization
- Just add a decorator to your compute intensive function

Mandelbrot in Numba

```
from numba import autojit
@autojit # <--- All we need to add
def mandel(x, y, max_iters):
    i = 0
    c = complex(x,y)
    z = 0.0j
    for i in range(max_iters):
        z = z*z + c
        if (z.real*z.real + z.imag*z.imag) >= 4:
            return i
    return 255
```

- Over **170x** speedup

Need More Speed? Use NumbaPro “CUDA Python”

- Our commerical product enables parallel compute on GPU.

```
from numbapro import cuda, uint8, f8, uint32
# use CUDA jit
@cuda.jit(argtypes=[uint8[:,:], f8, f8, f8, f8, uint32])
def mandel_kernel(image, min_x, max_x,
                  min_y, max_y, iters):
    height = image.shape[0]
    width = image.shape[1]
    pixel_size_x = (max_x - min_x) / width
    pixel_size_y = (max_y - min_y) / height
    # access thread indices
    x = cuda.threadIdx.x + \
        cuda.blockIdx.x * cuda.blockDim.x
    y = cuda.threadIdx.y + \
        cuda.blockIdx.y * cuda.blockDim.y
    # truncated ...
```

- **1255x** faster on K20

Lab 1: Saxpy in “CUDA Python”

- Implement saxpy in “CUDA Python”
- The lab is broken down into four small exercises
- We will provide guidelines and hints along the way
- lab1/saxpy.py

Exercise 1

Host -> Device

- `d_ary = cuda.to_device(ary)`
 - `cudaMalloc(size);`
 - `cudaMemcpy(devary, hstary, size, cudaMemcpyHostToDevice);`

Host -> Device (allocate only, no copy)

- `d_ary = cuda.to_device(ary, copy=False)`
 - `cudaMalloc(size);`

Exercise 2

Kernel Launch

- griddim: tuple of 1-2 ints
- blockdim: tuple of 1-3 ints
 - 'dim3 griddim, blockdim;
- `a_kernel[griddim, blockddim](arg0, arg1)`
 - `a_kernel<<<griddim, blockdim>>>(arg0, arg1);`

Exercise 3

Device -> Host

- `d_ary.to_host()`
 - `cudaMemcpy(hstary, devary, size, cudaMemcpyHostToDevice);`

Exercise 4

Inside the Kernel

- `cuda.threadIdx`, `cuda.blockIdx`, `cuda.blockDim`
 - `threadIdx`, `blockIdx`, `blockDim`

```
i = cuda.threadIdx.x + cuda.blockIdx.x * cuda.blockDim.x
```

“CUDA Python” Too Low-Level?

- “CUDA Python” API may be too close to CUDA-C
- We want simpler API
- An API that is closer to array-expression

NumbaPro Compute Unit (CU) API

- Heterogeneous parallel programming for GPU, CPU, ...
- Execute kernels asynchronously

A Saxpy Example

```
def product(tid, A, B, Prod):  
    Prod[tid] = A[tid] * B[tid]
```

```
def sum(tid, A, B, Sum):  
    Sum[tid] = A[tid] + B[tid]
```

```
cu = CU('gpu') # or CU('cpu')  
... # prepare data  
cu.enqueue(product, ntid=dProd.size,  
            args=(dA, dB, dProd))  
cu.enqueue(sum,      ntid=dSum.size,  
            args=(dProd, dC, dSum))  
... # do something while waiting?  
cu.wait()  
  
cu.close() # destroy the compute unit
```

Lab 2: A “Monty” Carlo Option Pricer in CU API

- Implement a monte carlo pricer using the CU API
- Guidelines and hints will be provided as we go
- lab2/pricer_cu.py

Step 0

Instantiate a Compute Unit

```
cu = CU('gpu')    # or 'cpu' for multicore
```


Step 1

Prepare Data Memory

- read only
 - `d_ary = cu.input(ary)`
- write only
 - `d_ary = cu.output(ary)`
- read+write
 - `d_ary = cu.inout(ary)`
- scratchpad
 - `d_ary = cu.scratch(shape=arraylen, dtype=np.float32)`
 - `d_ary = cu.scratch_like(ary)`

Exercise 1

```
d_noises = # fill in the RHS  
# Hints: length of array is n
```

Step 2

Enqueue kernels

- `cu.enqueue(kernel, ntid=number_of_threads, args=(arg0, arg1))`
 - `tid` (1st argument of the kernel) is automatically populated
- Kernels run asynchronously

Exercise 2

- Enqueue the “step” kernel

Step 3

Wait for the kernel to complete

- `cu.wait()`

Step 4

Fill in the kernel

Exercise 3

- Use the Numpy version as a reference.

A Numpy Implementation

```
import numpy as np
from math import sqrt, exp
from timeit import default_timer as timer

def step(dt, prices, c0, c1, noises):
    return prices * np.exp(c0 * dt + c1 * noises)

def monte_carlo_pricer(paths, dt, interest, volatility):
    c0 = interest - 0.5 * volatility ** 2
    c1 = volatility * np.sqrt(dt)

    for j in xrange(1, paths.shape[1]):
        prices = paths[:, j - 1]
        noises = np.random.normal(0., 1., prices.size)
        paths[:, j] = step(dt, prices, c0, c1, noises)

if __name__ == '__main__':
    from driver import driver
    driver(monte_carlo_pricer)
```

Expected Result

The result should be close to the following numbers:

```
StockPrice 22.6403957688  
StandardError 0.000434370525451  
PaidOff 1.14039936311  
OptionPrice 1.04921806448
```

Numpy implementation

- 19.74 MStep per second

NumbaPro CU + GeForce GT 650M

- 101.78 MStep per second
- **5x** speedup

NumbaPro CU + Tesla C2075

- 188.84 MStep per second
- **9.5x** speedup

Questions?